

Homework 2

Shreyas Bidwai

October 4, 2017

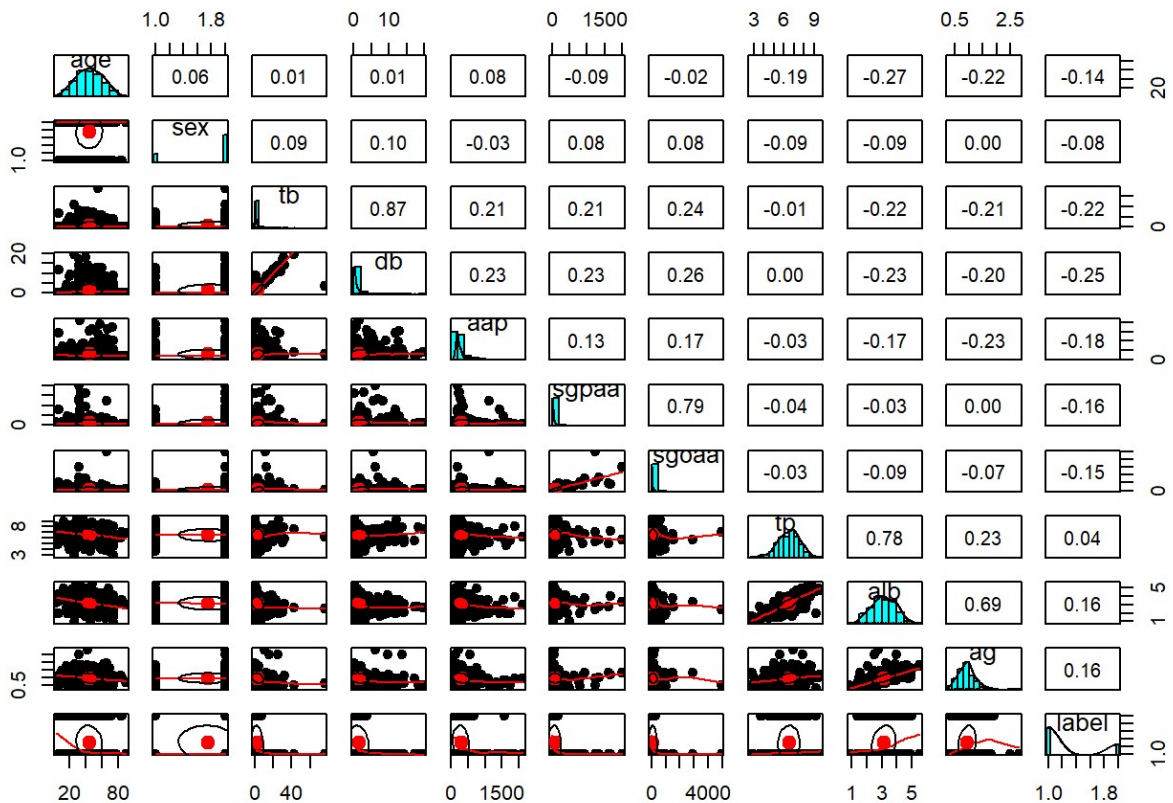
```
library(rpart)
library(caret)
library(rpart.plot)
library(ROCR)

ilpd=read.csv("/ILPD.csv", header = T, sep=",")

set.seed(100)

index<- sample(1:nrow(ilpd), size = 0.6*nrow(ilpd))
train<- ilpd[index, ]
test <- ilpd[-index, ]

library(psych)
pairs.panels(ilpd, pch=19)
```



a)

i) Strongest correlated pair :- db and tb

ii) Weakest correlated pair :- tp and db, sex and ag, sgpa and ag

iii) Most negatively correlated :- age and alb

iv) Variables that appear to follow a Gaussian distribution :- age, tp, alb, ag

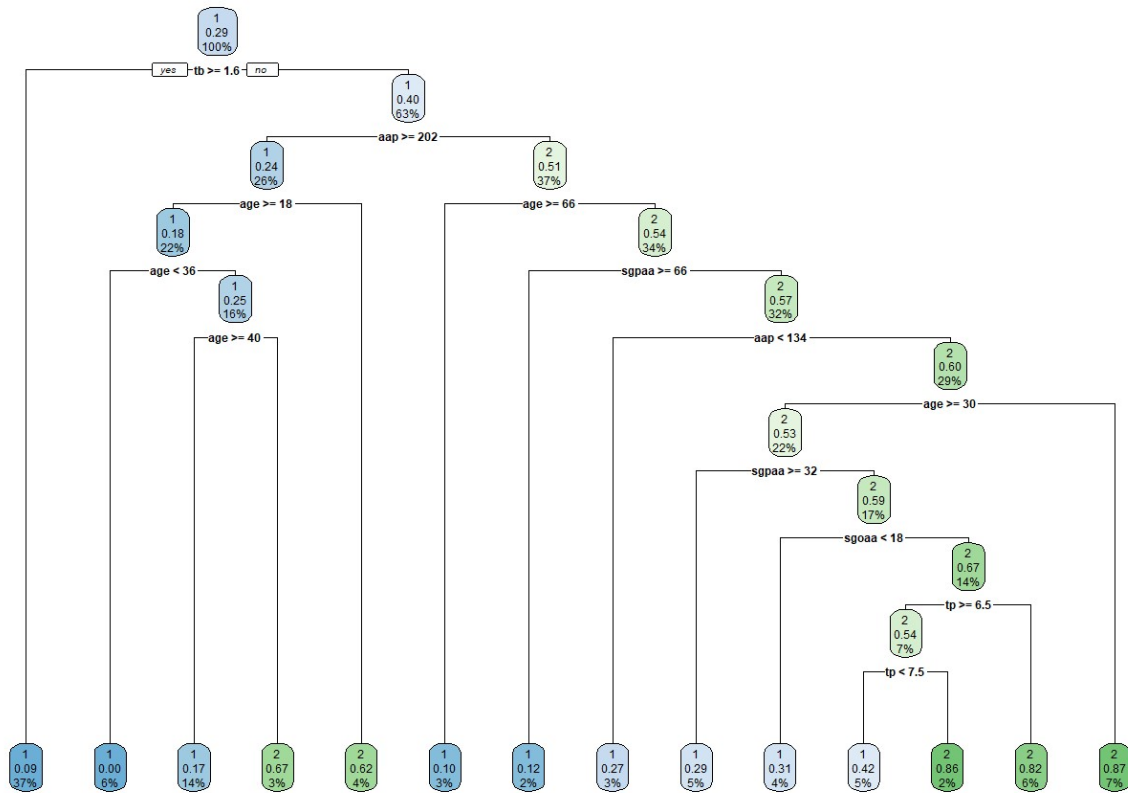
b)

Yes, I think normalising or scaling the attributes will help the classification task. Because, normalising the attributes will result in data which are similar to each other. It will also help in providing a better correlation and there is no point in normalising the data that are similar to each other.

Attributes with varied range of values that should be normalised are:- Age, tp, alb, ag

c)

```
model <- rpart(label~ ., method = "class", data = train)
rpart.plot(model)
```



```

pred <- predict(model, test, type = "class")
confusionMatrix(pred, test[,11], positive = "1")

```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction   1   2
##           1 145  51
##           2  22  16
##
##           Accuracy : 0.688
##           95% CI : (0.6244, 0.7468)
##           No Information Rate : 0.7137
##           P-Value [Acc > NIR] : 0.826721
##
##           Kappa : 0.123
##           Mcnemar's Test P-Value : 0.001049
##
##           Sensitivity : 0.8683
##           Specificity : 0.2388
##           Pos Pred Value : 0.7398
##           Neg Pred Value : 0.4211
##           Prevalence : 0.7137
##           Detection Rate : 0.6197
##           Detection Prevalence : 0.8376
##           Balanced Accuracy : 0.5535
##
##           'Positive' Class : 1
##
```

Accuracy: 68.8%

TPR(Sensitivity) : 0.8683


TNR(Specificity) : 0.2388

PPV(Pos Pred Value) : 0.7398

d)

-Prune

```
plotcp(model)
```



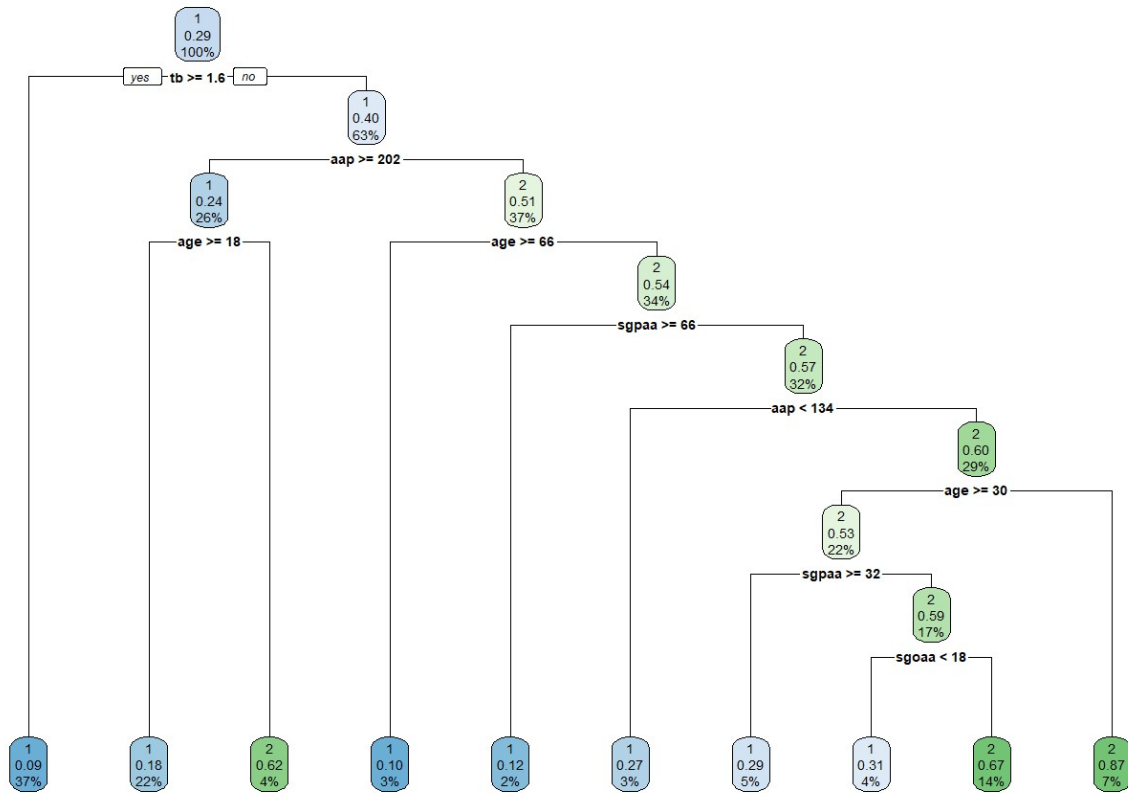
```
printcp(model)
```

```
##
## Classification tree:
## rpart(formula = label ~ ., data = train, method = "class")
##
## Variables actually used in tree construction:
## [1] aap  age  sgoaa sgpaa tb    tp
##
## Root node error: 100/349 = 0.28653
##
## n= 349
##
##      CP nsplit rel error xerror   xstd
## 1 0.033333     0    1.00  1.00 0.084467
## 2 0.030000     8    0.67  1.25 0.089571
## 3 0.015000     9    0.64  1.18 0.088376
## 4 0.010000    13    0.58  1.22 0.089080
```

```
model.pruned <- prune(model, cp = 0.021)
pred.pruned <- predict(model.pruned, test, type = "class")
confusionMatrix(pred.pruned, test[, 11], positive = "1")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  1    2
##           1 142  46
##           2  25  21
##
##           Accuracy : 0.6966
##           95% CI : (0.6333, 0.7548)
##    No Information Rate : 0.7137
##    P-Value [Acc > NIR] : 0.74428
##
##           Kappa : 0.1807
##  McNemar's Test P-Value : 0.01762
##
##           Sensitivity : 0.8503
##           Specificity : 0.3134
##           Pos Pred Value : 0.7553
##           Neg Pred Value : 0.4565
##           Prevalence : 0.7137
##           Detection Rate : 0.6068
##    Detection Prevalence : 0.8034
##           Balanced Accuracy : 0.5819
##
##           'Positive' Class : 1
##
```

```
rpart.plot(model.pruned)
```

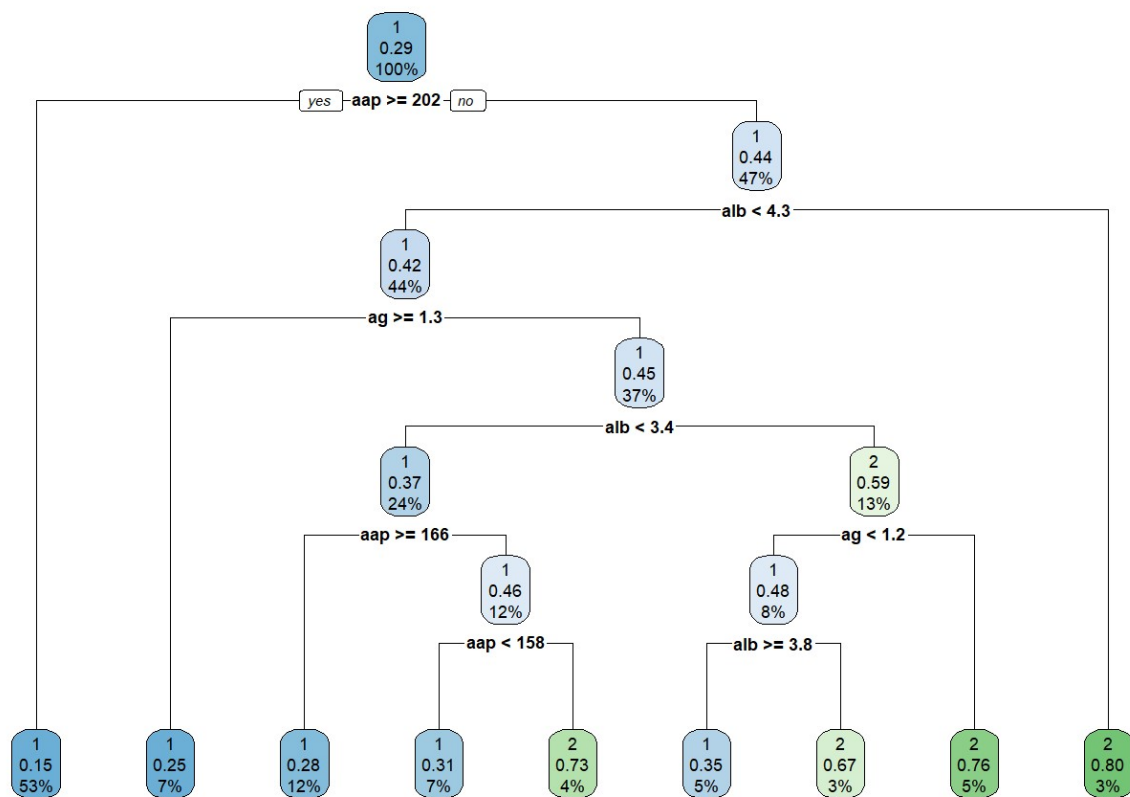


Accuracy of the model is 68.8% and after changing the values of cp I got a better accuracy of 69.66%.

Thus pruned tree has more accuracy because of lesser complexity, as there are less number of nodes to traverse in decision tree.

e)Build a new model

```
new_model<- rpart(label ~ alb+ag+aap, method = "class", data = train)
rpart.plot(new_model)
```



```

new_pred <- predict(new_model,test,type = "class")
confusionMatrix(new_pred,test[,11], positive = "1")

```



```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction  1    2
##           1 152  49
##           2   15  18
##
##           Accuracy : 0.7265
##           95% CI : (0.6646, 0.7825)
##           No Information Rate : 0.7137
##           P-Value [Acc > NIR] : 0.3623
##
##           Kappa : 0.2109
##           McNemar's Test P-Value : 3.707e-05
##
##           Sensitivity : 0.9102
##           Specificity : 0.2687
##           Pos Pred Value : 0.7562
##           Neg Pred Value : 0.5455
##           Prevalence : 0.7137
##           Detection Rate : 0.6496
##           Detection Prevalence : 0.8590
##           Balanced Accuracy : 0.5894
##
##           'Positive' Class : 1
##
```

```
summary(ilpd)
```

```
##      age      sex      tb      db
## Min.   : 4.00  Female:142  Min.   : 0.400  Min.   : 0.100
## 1st Qu.:33.00  Male  :441  1st Qu.: 0.800  1st Qu.: 0.200
## Median :45.00           Median : 1.000  Median : 0.300
## Mean   :44.75           Mean   : 3.299  Mean   : 1.486
## 3rd Qu.:58.00           3rd Qu.: 2.600  3rd Qu.: 1.300
## Max.   :90.00           Max.   :75.000  Max.   :19.700
##      aap      sgpa      sgoaa      tp
## Min.   : 63.0  Min.   : 10.00  Min.   : 10.0  Min.   :2.700
## 1st Qu.:175.5  1st Qu.: 23.00  1st Qu.: 25.0  1st Qu.:5.800
## Median :208.0  Median : 35.00  Median : 42.0  Median :6.600
## Mean   :290.6  Mean   : 80.71  Mean   :109.9  Mean   :6.483
## 3rd Qu.:298.0  3rd Qu.: 60.50  3rd Qu.: 87.0  3rd Qu.:7.200
## Max.   :2110.0  Max.   :2000.00  Max.   :4929.0  Max.   :9.600
##      alb      ag      label
## Min.   :0.900  Min.   :0.300  Min.   :1.000
## 1st Qu.:2.600  1st Qu.:0.700  1st Qu.:1.000
## Median :3.100  Median :0.940  Median :1.000
## Mean   :3.142  Mean   :0.947  Mean   :1.286
## 3rd Qu.:3.800  3rd Qu.:1.100  3rd Qu.:2.000
## Max.   :5.500  Max.   :2.800  Max.   :2.000
```

Accuracy :- 72.65%

TPR(Sensitivity) : 0.9102

TNR(Specificity) : 0.2687


PPV(Pos Pred Value) : 0.7562

f)

(i) a ROC curve using the ROCR package.

ROC for first model

```
pred.roc <- predict(model,newdata=test,type="prob")[,2]
f.pred <- prediction(pred.roc,test$label)
f.perf <- performance(f.pred, "tpr", "fpr")
plot(f.perf, colorize=T, lwd=3, main = "ROC")
abline(0,1)
```




```
auc<-performance(f.pred,measure = "auc")  
auc@y.values[[1]]
```

```
## [1] 0.6675753
```

ROC for new model

```
pred.roc1 <- predict(new_model,newdata=test,type="prob")[,2]  
f.pred1 <- prediction(pred.roc1,test$label)  
f.perf1 <- performance(f.pred1, "tpr", "fpr")  
plot(f.perf1, colorize=T, lwd=3, main = "ROC")  
abline(0,1)
```



```
auc<-performance(f.pred1,measure = "auc")  
auc@y.values[[1]]
```

```
## [1] 0.6455
```

ii) AUC for model: 0.6675

AUC for newmodel: 0.6455

iii) Previous Model is better because the auc is closer to 1 and greater than as compared to new model.